

```
import os
import numpy as np
import torch
import torch.nn as nn
from torch.nn.utils import clip_grad_norm_
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms
from tqdm.notebook import tqdm
from torchvision.transforms.functional import InterpolationMode

import matplotlib.pyplot as plt
from PIL import Image

import shutil

seed = 142

np.random.seed(seed)
torch.manual_seed(seed)
torch.cuda.manual_seed(seed)
torch.backends.cudnn.deterministic = True

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

cuda:0

class ImageDataset(Dataset):
    def __init__(self, data_dir, transforms=None):
        monet_dir = os.path.join(data_dir, 'monet_jpg')
        photo_dir = os.path.join(data_dir, 'photo_jpg')

        self.files_monet = [os.path.join(monet_dir, name) for name in sorted(os.listdir(monet_dir))]
        self.files_photo = [os.path.join(photo_dir, name) for name in sorted(os.listdir(photo_dir))]

        self.transforms = transforms

    def __len__(self):
        # we know that len(files_monet) = 300 < 7038 = len(files_photo)
        return len(self.files_monet)

    def __getitem__(self, index):
        # we will use only 300 (=len(files_monet)) photos during training
        # randomly picking them from the first 300 photos
        random_index = np.random.randint(0, len(self.files_monet))
        file_monet = self.files_monet[index]
        file_photo = self.files_photo[random_index]
```

```

image_monet = Image.open(file_monet)
image_photo = Image.open(file_photo)

if self.transforms is not None:
    image_monet = self.transforms(image_monet)
    image_photo = self.transforms(image_photo)

return image_monet, image_photo

data_dir = '/content/drive/MyDrive/Colab Notebooks/kaggle/data'
batch_size = 5

transforms_ = transforms.Compose([
    #transforms.Resize((256, 256)), # photos already have the same size
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomVerticalFlip(p=0.3),
    transforms.ToTensor(),
    transforms.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5)),
])
dataloader = DataLoader(
    ImageDataset(data_dir, transforms=transforms_),
    batch_size=batch_size,
    shuffle=True,
    num_workers=2,
)
def unnorm(img, mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]):
    for t, m, s in zip(img, mean, std):
        t.mul_(s).add_(m)

    return img

```

▼ Build CycleGAN

▼ Auxiliary blocks

```

class ConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=3, stride=1, padding=0,
                 transpose=False, use_leaky=True, use_dropout=False, normalize=True):

        super(ConvBlock, self).__init__()
        self.block = []

```

```
if transpose:
    self.block += [nn.ConvTranspose2d(in_channels, out_channels, kernel_size,
                                    stride, padding, output_padding=1)]
else:
    self.block += [nn.Conv2d(in_channels, out_channels, kernel_size,
                           stride, padding, bias=True)]

if normalize:
    self.block += [nn.InstanceNorm2d(out_channels)]

if use_dropout:
    self.block += [nn.Dropout(0.5)]

if use_leaky:
    self.block += [nn.LeakyReLU(negative_slope=0.2, inplace=True)]
else:
    self.block += [nn.ReLU(inplace=True)]

self.block = nn.Sequential(*self.block)

def forward(self, x):
    return self.block(x)

class ResidualBlock(nn.Module):
    def __init__(self, channels):
        super(ResidualBlock, self).__init__()
        self.block = nn.Sequential(
            nn.ReflectionPad2d(1),
            ConvBlock(in_channels=channels, out_channels=channels,
                      kernel_size=3, use_leaky=False, use_dropout=True),
            nn.ReflectionPad2d(1),
            nn.Conv2d(in_channels=channels, out_channels=channels, kernel_size=3),
            nn.InstanceNorm2d(channels)
        )

        def forward(self, x):
            return x + self.block(x)

class ConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=3, stride=1, padding=1,
                 transpose=False, use_leaky=True, use_dropout=False, normalize=True):
        super(ConvBlock, self).__init__()
        self.block = []

        if transpose:
            self.block += [nn.ConvTranspose2d(in_channels, out_channels, kernel_size,
```

```
        stride, padding, output_padding=1)]  
    else:  
        self.block += [nn.Conv2d(in_channels, out_channels, kernel_size,  
                               stride, padding, bias=True)]  
  
    if normalize:  
        self.block += [nn.InstanceNorm2d(out_channels, affine=True)]  
  
    if use_dropout:  
        self.block += [nn.Dropout(0.5)]  
  
    if use_leaky:  
        self.block += [nn.LeakyReLU(negative_slope=0.2, inplace=True)]  
    else:  
        self.block += [nn.ReLU(inplace=True)]  
  
    self.block = nn.Sequential(*self.block)  
  
def forward(self, x):  
    return self.block(x)  
  
class ResidualBlock(nn.Module):  
    def __init__(self, channels, kernel_size=3):  
        super(ResidualBlock, self).__init__()  
  
        #doesn't change shape of input  
        self.block = nn.Sequential(  
            nn.ReflectionPad2d(int((kernel_size - 1) / 2)),  
            ConvBlock(in_channels=channels, out_channels=channels, padding=0,  
                      kernel_size=kernel_size, use_leaky=False, use_dropout=True),  
            nn.ReflectionPad2d(int((kernel_size - 1) / 2)),  
            nn.Conv2d(in_channels=channels, out_channels=channels, kernel_size=kernel_size, pa  
            nn.InstanceNorm2d(channels, affine=True)  
    )  
  
    def forward(self, x):  
        return x + self.block(x)  
  
class transformer_(nn.Module):  
    def __init__(self, transformer_kernel_size=9):  
        super(transformer_, self).__init__()  
        self.model = nn.AvgPool2d(transformer_kernel_size, stride=1, padding=transformer_kernel_  
  
    def forward(self, x):  
        return self.model(x)
```

▼ Define Generator

```
class Generator(nn.Module):
    def __init__(self, in_channels, out_channels, num_residual_blocks=9):
        super(Generator, self).__init__()

        ''' Encoder '''
        # Initial layer: 3*256*256 -> 64*256*256
        self.initial = [
            nn.ReflectionPad2d(in_channels),
            ConvBlock(in_channels=in_channels, out_channels=64,
                      kernel_size=2*in_channels+1, use_leaky=False),
        ]
        self.initial = nn.Sequential(*self.initial)

        # Downsampling: 64*256*256 -> 128*128*128 -> 256*64*64
        self.down = [
            ConvBlock(in_channels=64, out_channels=128, kernel_size=3,
                      stride=2, padding=1, use_leaky=False),
            ConvBlock(in_channels=128, out_channels=256, kernel_size=3,
                      stride=2, padding=1, use_leaky=False),
        ]
        self.down = nn.Sequential(*self.down)

        """ Transformer """
        # ResNet: 256*64*64 -> 256*64*64
        self.transform = [ResidualBlock(256) for _ in range(num_residual_blocks)]
        self.transform = nn.Sequential(*self.transform)

        """ Decoder """
        # Upsampling: 256*64*64 -> 128*128*128 -> 64*256*256
        self.up = [
            ConvBlock(in_channels=256, out_channels=128, kernel_size=3, stride=2,
                      padding=1, transpose=True, use_leaky=False),
            ConvBlock(in_channels=128, out_channels=64, kernel_size=3, stride=2,
                      padding=1, transpose=True, use_leaky=False),
        ]
        self.up = nn.Sequential(*self.up)

        # Out layer: 64*256*256 -> 3*256*256
        self.out = nn.Sequential(
            nn.ReflectionPad2d(out_channels),
```

```

        nn.Conv2d(in_channels=64, out_channels=out_channels, kernel_size=2*out_channels+1
        nn.Tanh()
    )

def forward(self, x):
    x = self.down(self.initial(x))
    x = self.transform(x)
    x = self.out(self.up(x))
    return x

class encoder_(nn.Module):

    def __init__(self,in_channels=3,n_filter_generator=32,pad_type=0):
        super(encoder_, self).__init__()

        self.model=nn.Sequential(
            nn.InstanceNorm2d(in_channels,affine=True),
            #3*256*256 - 3*286*286
            nn.ReflectionPad2d(15),
            #3*286*286 - 32*284*284
            ConvBlock(in_channels=in_channels, out_channels=n_filter_generator, kernel_size=3
                      stride=1, padding=pad_type, normalize=True,use_leaky=False),
            #32*284*284 - 32*141*141
            ConvBlock(in_channels=n_filter_generator, out_channels=n_filter_generator, kernel_
                      stride=2, padding=pad_type, normalize=True,use_leaky=False),
            #32*141*141 - 64*70*70
            ConvBlock(in_channels=n_filter_generator, out_channels=n_filter_generator*2, kern
                      stride=2, padding=pad_type, normalize=True,use_leaky=False),
            #64*70*70 - 128*34*34
            ConvBlock(in_channels=n_filter_generator*2, out_channels=n_filter_generator*4, ke
                      stride=2, padding=pad_type, normalize=True,use_leaky=False),
            #128*34*34 - 256*32*32
            ConvBlock(in_channels=n_filter_generator*4, out_channels=n_filter_generator*8, ke
                      stride=1, padding=pad_type, normalize=True,use_leaky=False)
        )

    def forward(self, x):
        return self.model(x)

class decoder_(nn.Module):

    """decoder model following https://arxiv.org/pdf/1807.10201.pdf
    Returns: decoder model
    """
    def __init__(self, input_shape,n_filter_generator=32):
        super(decoder_, self).__init__()

```

```

num_kernels = n_filter_generator * 8
self.model=[]
for i in range(9):
    self.model.append(ResidualBlock(num_kernels))

#i=0: 256*32*32 - 256*32*32 - 256*32*32
#i=1: 256*32*32 - 256*64*64 - 128*64*64
#i=2: 128*64*64 - 128*128*128 - 64*128*128
#i=3: 64*128*128 - 64*256*256 - 32*256*256
in_channels=[num_kernels,8*n_filter_generator,4*n_filter_generator,2*n_filter_generator]
for i in range(4):
    self.model.append(transforms.Resize(size=(input_shape[2]*2***(i),input_shape[3]*2***(i))))
    self.model.append(nn.Conv2d(in_channels=in_channels[i], out_channels=n_filter_generator, kernel_size=3, stride=1, padding=1, bias=True))
    self.model.append(nn.InstanceNorm2d(n_filter_generator * 2 ** (3 - i), affine=True))
    self.model.append(nn.ReLU(inplace=True))

self.model.append(nn.ReflectionPad2d(3))

#32*256*256 - 3*256*256
self.model.append(nn.Conv2d(in_channels=n_filter_generator, out_channels=3, kernel_size=3, stride=1, padding=0, bias=True))
self.model.append(nn.Tanh())
self.model = nn.Sequential(*self.model)
#x = CenterLayer()(x)

def forward(self,x):
    return self.model(x)##*2-1

```

▼ Define Discriminator

```

class discriminator_(nn.Module):
    def __init__(self, in_channels=3,pad_type='valid',n_filter_discriminator=64):#ASA code uses
        super(discriminator_, self).__init__()

    #3*256*256 - 128*128*128 -> 1*128*128
    self.conv0=ConvBlock(in_channels=in_channels, out_channels=n_filter_discriminator*2,
                        stride=2, padding=2, normalize=True)
    self.conv0_pred=nn.Conv2d(in_channels=n_filter_discriminator*2, out_channels=1, kernel_size=1, stride=1, padding=2, bias=True)

    #128*128*128 - 128*64*64 -> 1*65*65
    self.conv1=ConvBlock(in_channels=n_filter_discriminator*2, out_channels=n_filter_discriminator*4,
                        stride=2, padding=2, normalize=True)
    self.conv1_pred=nn.Conv2d(in_channels=n_filter_discriminator*4, out_channels=1, kernel_size=1, stride=1, padding=5, bias=True)

```

```
#128*64*64 - 256*32*32
self.conv2=ConvBlock(in_channels=n_filter_discriminator*2, out_channels=n_filter_disc
                    stride=2, padding=2, normalize=True)

#256*32*32 - 512*16*16 -> 1*17*17
self.conv3=ConvBlock(in_channels=n_filter_discriminator*4, out_channels=n_filter_disc
                    stride=2, padding=2, normalize=True)
self.conv3_pred=nn.Conv2d(in_channels=n_filter_discriminator*8, out_channels=1, kernel_
                    stride=1, padding=5, bias=True)

#512*16*16 - 512*8*8
self.conv4=ConvBlock(in_channels=n_filter_discriminator*8, out_channels=n_filter_disc
                    stride=2, padding=2, normalize=True)

#512*4*4 - 1024*4*4 -> 1*5*5
self.conv5=ConvBlock(in_channels=n_filter_discriminator*8, out_channels=n_filter_disc
                    stride=2, padding=2, normalize=True)
self.conv5_pred=nn.Conv2d(in_channels=n_filter_discriminator*16, out_channels=1, kernel_
                    stride=1, padding=3, bias=True)

#1024*1*1 - 1024*2*2 -> 1*2*2
self.conv6=ConvBlock(in_channels=n_filter_discriminator*16, out_channels=n_filter_disc
                    stride=2, padding=0, normalize=True)
self.conv6_pred=nn.Conv2d(in_channels=n_filter_discriminator*16, out_channels=1, kernel_
                    stride=1, padding=1, bias=True)

def forward(self, x):

    h0=self.conv0(x)
    h0_pred=self.conv0_pred(h0)

    h1=self.conv1(h0)
    h1_pred=self.conv1_pred(h1)

    h2=self.conv2(h1)

    h3=self.conv3(h2)
    h3_pred=self.conv3_pred(h3)

    h4=self.conv4(h3)

    h5=self.conv5(h4)
    h5_pred=self.conv5_pred(h5)

    h6=self.conv6(h5)
    h6_pred=self.conv6_pred(h6)

    #print((h0.size(),h1.size(),h2.size(),h3.size(),h4.size(),h5.size(),h6.size()))
```

```
#print((h0_pred.size(),h1_pred.size(),h3_pred.size(),h5_pred.size(),h6_pred.size()))
return [h0_pred, h1_pred, h3_pred, h5_pred, h6_pred]
```

```
class Discriminator(nn.Module):
    def __init__(self, in_channels):
        super(Discriminator, self).__init__()

        self.model = nn.Sequential(
            # 3*256*256 -> 64*128*128
            ConvBlock(in_channels=in_channels, out_channels=64, kernel_size=4,
                      stride=2, padding=1, normalize=False),

            # 64*128*128 -> 128*64*64
            ConvBlock(in_channels=64, out_channels=128, kernel_size=4, stride=2, padding=1),

            # 128*64*64 -> 256*32*32
            ConvBlock(in_channels=128, out_channels=256, kernel_size=4, stride=2, padding=1),

            # 256*32*32 -> 512*31*31
            ConvBlock(in_channels=256, out_channels=512, kernel_size=4, stride=1, padding=1),

            # 512*31*31 -> 1*30*30
            nn.Conv2d(in_channels=512, out_channels=1, kernel_size=4, stride=1, padding=1),
        )

    def forward(self, x):
        return self.model(x)
```

▼ Training tools

▼ Model initialization

```
generator_monet2photo = Generator(in_channels=3, out_channels=3, num_residual_blocks=9).to(device)
generator_photo2monet = Generator(in_channels=3, out_channels=3, num_residual_blocks=9).to(device)

discriminator_monet = Discriminator(in_channels=3).to(device)
discriminator_photo = Discriminator(in_channels=3).to(device)

generator_monet2photo = nn.Sequential(encoder_(),decoder_(input_shape=(batch_size,256,32,32)))
generator_photo2monet = nn.Sequential(encoder_(),decoder_(input_shape=(batch_size,256,32,32)))
```

```
discriminator_monet = Discriminator(in_channels=3).to(device)
discriminator_photo = Discriminator(in_channels=3).to(device)
```

▼ Define Losses

```
criterion_GAN = nn.MSELoss()
criterion_cycle = nn.L1Loss()
criterion_identity = nn.L1Loss()
```

▼ Define Optimizers

```
lr = 2e-4
b1 = 0.5
b2 = 0.999

optim_generators = torch.optim.Adam(
    list(generator_monet2photo.parameters()) + list(generator_photo2monet.parameters()),
    lr=lr, betas=(b1, b2)
)

optim_discriminators = torch.optim.Adam(
    list(discriminator_monet.parameters()) + list(discriminator_photo.parameters()),
    lr=lr, betas=(b1, b2)
)
```

▼ Define Learning rate schedulers

```
num_epochs = 80
decay_epoch = 25

lr_sched_step = lambda epoch: 1 - max(0, epoch - decay_epoch) / (num_epochs - decay_epoch)

lr_sched_generators = torch.optim.lr_scheduler.LambdaLR(optim_generators, lr_lambda=lr_sched_step)
lr_sched_discriminators = torch.optim.lr_scheduler.LambdaLR(optim_discriminators, lr_lambda=1 - lr_sched_step)
```

▼ Auxiliary tools

```
class History:
    def __init__(self):
        self.generators_loss = []
        self.discriminators_loss = []
```

```

def update(self, gen_loss, discr_loss):
    self.generators_loss.append(gen_loss)
    self.discriminators_loss.append(discr_loss)

def show(self, title='Losses'):
    fig = plt.figure(figsize=(20, 8))
    plt.title(title)
    plt.plot(self.generators_loss, 'o-', color='r',
              linewidth=2, markersize=3, label='Generators Loss')
    plt.plot(self.discriminators_loss, 'o-', color='b',
              linewidth=2, markersize=3, label='Discriminators Loss')
    plt.legend(loc='best')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.grid(True)
    plt.show()

class Buffer:
    def __init__(self, max_images=50):
        self.max_images = max_images
        self.images = []

    def update(self, images):
        images = images.detach().cpu().data.numpy()
        for image in images:
            if len(self.images) < self.max_images:
                self.images.append(image)
            else:
                if np.random.rand() > 0.5:
                    index = np.random.randint(0, self.max_images)
                    self.images[index] = image

    def sample(self, num_images):
        samples = np.array([self.images[np.random.randint(0, len(self.images))]]
                           for _ in range(num_images)])
        return torch.tensor(samples)

def update_req_grad(models, requires_grad=True):
    for model in models:
        for param in model.parameters():
            param.requires_grad = requires_grad

```

▼ Training CycleGAN

```

history = History()
buffer_monet = Buffer()
buffer_photo = Buffer()

for epoch in range(num_epochs):

```

```
avg_generators_loss = 0
avg_discriminators_loss = 0

for i, (real_monet, real_photo) in enumerate(tqdm(dataloader, leave=False, total=len(data)):
    real_monet, real_photo = real_monet.to(device), real_photo.to(device)

    """ Train Generators """
    # switching models parameters so that only generators are trained
    update_req_grad([generator_monet2photo, generator_photo2monet], True)
    update_req_grad([discriminator_monet, discriminator_photo], False)

    # zero the parameters gradients
    optim_generators.zero_grad()

    # forward-pass
    fake_photo = generator_monet2photo(real_monet)
    fake_monet = generator_photo2monet(real_photo)

    cycle_photo = generator_monet2photo(fake_monet)
    cycle_monet = generator_photo2monet(fake_photo)

    identity_photo = generator_monet2photo(real_photo)
    identity_monet = generator_photo2monet(real_monet)

    # update photos that are used to feed up discriminators
    buffer_photo.update(fake_photo)
    buffer_monet.update(fake_monet)

    # discriminators outputs that are used in adversarial loss
    discriminator_outputs_photo = discriminator_photo(fake_photo)
    discriminator_outputs_monet = discriminator_monet(fake_monet)

    # labels that are used as ground truth
    labels_real = torch.ones(discriminator_outputs_monet.size()).to(device)
    labels_fake = torch.zeros(discriminator_outputs_monet.size()).to(device)

    # adversarial loss - enforces that the generated output be of the appropriate domain
    loss_GAN_monet2photo = criterion_GAN(discriminator_outputs_photo, labels_real)
    loss_GAN_photo2monet = criterion_GAN(discriminator_outputs_monet, labels_real)
    loss_GAN = (loss_GAN_monet2photo + loss_GAN_photo2monet) / 2

    # cycle consistency loss - enforces that the input and output are recognizably the same
    loss_cycle_photo = criterion_cycle(cycle_photo, real_photo)
    loss_cycle_monet = criterion_cycle(cycle_monet, real_monet)
    loss_cycle = (loss_cycle_photo + loss_cycle_monet) / 2

    # identity mapping loss - helps preserve the color of the input images
    loss_identity_photo = criterion_identity(identity_photo, real_photo)
    loss_identity_monet = criterion_identity(identity_monet, real_monet)
    loss_identity = (loss_identity_photo + loss_identity_monet) / 2

    # total loss
    avg_generators_loss += loss_GAN + loss_cycle + loss_identity
    avg_discriminators_loss += loss_GAN + loss_cycle
```

```
loss_generators_total = loss_GAN + 10 * loss_cycle + 5 * loss_identity

# backward-pass
loss_generators_total.backward()
optim_generators.step()

# limiting gradient norms - if they exceed 100, something went wrong
clip_grad_norm_(generator_photo2monet.parameters(), 100)
clip_grad_norm_(generator_monet2photo.parameters(), 100)

""" Train Discriminators """
# switching models parameters so that only discriminators are trained
update_req_grad([discriminator_monet, discriminator_photo], True)
update_req_grad([generator_monet2photo, generator_photo2monet], False)

# zero the parameters gradients
optim_discriminators.zero_grad()

# sample images from 50 stored
fake_photo = buffer_photo.sample(num_images=batch_size).to(device)
fake_monet = buffer_monet.sample(num_images=batch_size).to(device)

# making labels noisy for discriminators so that they don't prevail over generators
threshold = min(1, 0.85 + (1 - 0.85) * epoch / (num_epochs // 2))
noisy_labels_real = (torch.rand(discriminator_outputs_monet.size()) < threshold).float()

# forward-pass + losses
loss_real_photo = criterion_GAN(discriminator_photo(real_photo), noisy_labels_real)
loss_fake_photo = criterion_GAN(discriminator_photo(fake_photo.detach()), labels_fake)
loss_photo = (loss_real_photo + loss_fake_photo) / 2

loss_real_monet = criterion_GAN(discriminator_monet(real_monet), noisy_labels_real)
loss_fake_monet = criterion_GAN(discriminator_monet(fake_monet.detach()), labels_fake)
loss_monet = (loss_real_monet + loss_fake_monet) / 2

loss_discriminators_total = loss_monet + loss_photo

# backward-pass
loss_discriminators_total.backward()
optim_discriminators.step()

# clipping gradients to avoid gradients explosion
clip_grad_norm_(discriminator_monet.parameters(), 100)
clip_grad_norm_(discriminator_photo.parameters(), 100)

# updating intermediate results
avg_generators_loss += loss_generators_total.item()
avg_discriminators_loss += loss_discriminators_total.item()

# saving intermediate results
avg_generators_loss /= len(data_loader)
```

```
avg_generators_loss /= len(dataloader)
history.update(avg_generators_loss, avg_discriminators_loss)

# showing intermediate results
print("Epoch: %d/%d | Generators Loss: %.4f | Discriminators Loss: %.4f"
      % (epoch+1, num_epochs, avg_generators_loss, avg_discriminators_loss))

# showing generated images
if (epoch + 1) % 10 == 0:
    _, sample_real_photo = next(iter(dataloader))

    sample_fake_monet = generator_photo2monet(sample_real_photo.to(device)).detach().cpu()

    num_photos = min(batch_size, 5)
    plt.figure(figsize=(20, 8))
    for k in range(num_photos):
        plt.subplot(2, num_photos, k + 1)
        plt.imshow(unnorm(sample_real_photo[k]).permute(1, 2, 0))
        plt.title('Input photo')
        plt.axis('off')

        plt.subplot(2, num_photos, k + num_photos + 1)
        plt.imshow(unnorm(sample_fake_monet[k]).permute(1, 2, 0))
        plt.title('Output image')
        plt.axis('off')
    plt.show()

lr_sched_generators.step()
lr_sched_discriminators.step()
```





Output image



Output image



Output image



Output image



Output image

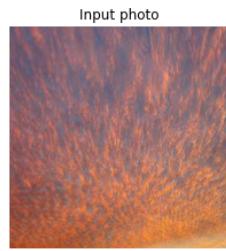
Epoch: 21/80	Generators Loss: 3.5662	Discriminators Loss: 0.2668
Epoch: 22/80	Generators Loss: 3.5757	Discriminators Loss: 0.2145
Epoch: 23/80	Generators Loss: 3.4235	Discriminators Loss: 0.2079
Epoch: 24/80	Generators Loss: 3.5432	Discriminators Loss: 0.2116
Epoch: 25/80	Generators Loss: 3.5553	Discriminators Loss: 0.2033
Epoch: 26/80	Generators Loss: 3.4547	Discriminators Loss: 0.2041
Epoch: 27/80	Generators Loss: 3.4947	Discriminators Loss: 0.1875
Epoch: 28/80	Generators Loss: 3.4394	Discriminators Loss: 0.1866
Epoch: 29/80	Generators Loss: 3.4496	Discriminators Loss: 0.1889
Epoch: 30/80	Generators Loss: 3.3074	Discriminators Loss: 0.1774



Input photo



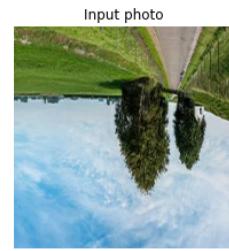
Input photo



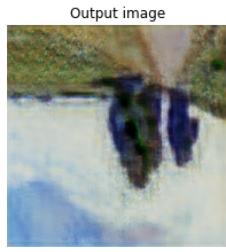
Input photo



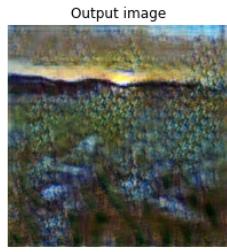
Input photo



Input photo



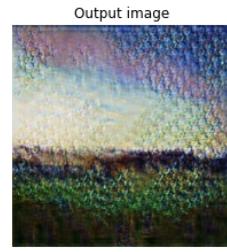
Output image



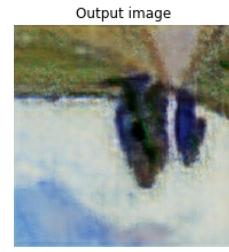
Output image



Output image



Output image



Output image

Epoch: 31/80	Generators Loss: 3.4079	Discriminators Loss: 0.1644
Epoch: 32/80	Generators Loss: 3.2475	Discriminators Loss: 0.1778
Epoch: 33/80	Generators Loss: 3.2962	Discriminators Loss: 0.1484
Epoch: 34/80	Generators Loss: 3.2317	Discriminators Loss: 0.1593
Epoch: 35/80	Generators Loss: 3.3812	Discriminators Loss: 0.6244
Epoch: 36/80	Generators Loss: 3.1057	Discriminators Loss: 0.3031
Epoch: 37/80	Generators Loss: 3.1574	Discriminators Loss: 0.2748
Epoch: 38/80	Generators Loss: 3.1272	Discriminators Loss: 0.2352
Epoch: 39/80	Generators Loss: 3.1643	Discriminators Loss: 0.2233
Epoch: 40/80	Generators Loss: 3.0134	Discriminators Loss: 0.1783



Input photo



Input photo



Input photo



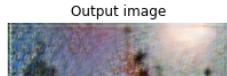
Input photo



Input photo



Output image



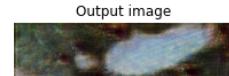
Output image



Output image



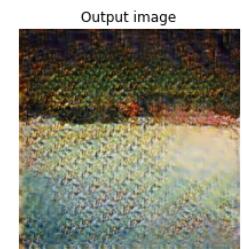
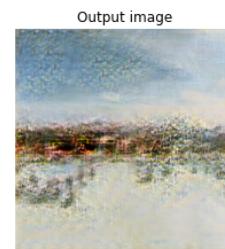
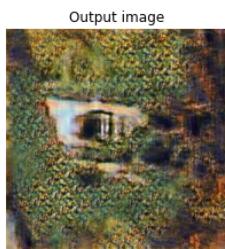
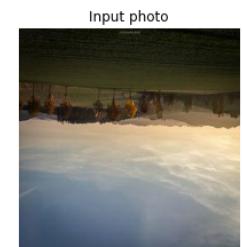
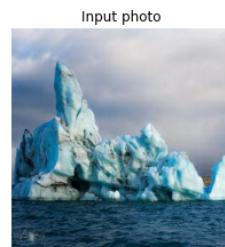
Output image



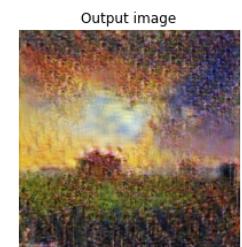
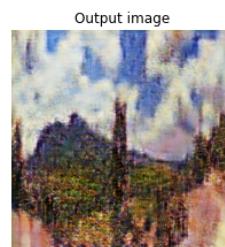
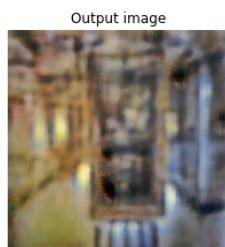
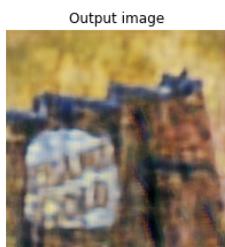
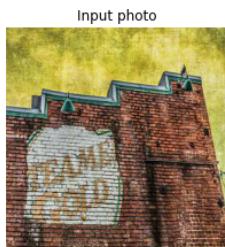
Output image



Epoch: 41/80	Generators Loss: 3.0460	Discriminators Loss: 0.1618
Epoch: 42/80	Generators Loss: 3.0042	Discriminators Loss: 0.1583
Epoch: 43/80	Generators Loss: 3.0283	Discriminators Loss: 0.1632
Epoch: 44/80	Generators Loss: 2.9768	Discriminators Loss: 0.1493
Epoch: 45/80	Generators Loss: 2.9740	Discriminators Loss: 0.1524
Epoch: 46/80	Generators Loss: 2.9119	Discriminators Loss: 0.1491
Epoch: 47/80	Generators Loss: 2.9590	Discriminators Loss: 0.1415
Epoch: 48/80	Generators Loss: 2.9422	Discriminators Loss: 0.1486
Epoch: 49/80	Generators Loss: 2.9141	Discriminators Loss: 0.1401
Epoch: 50/80	Generators Loss: 2.8690	Discriminators Loss: 0.1327

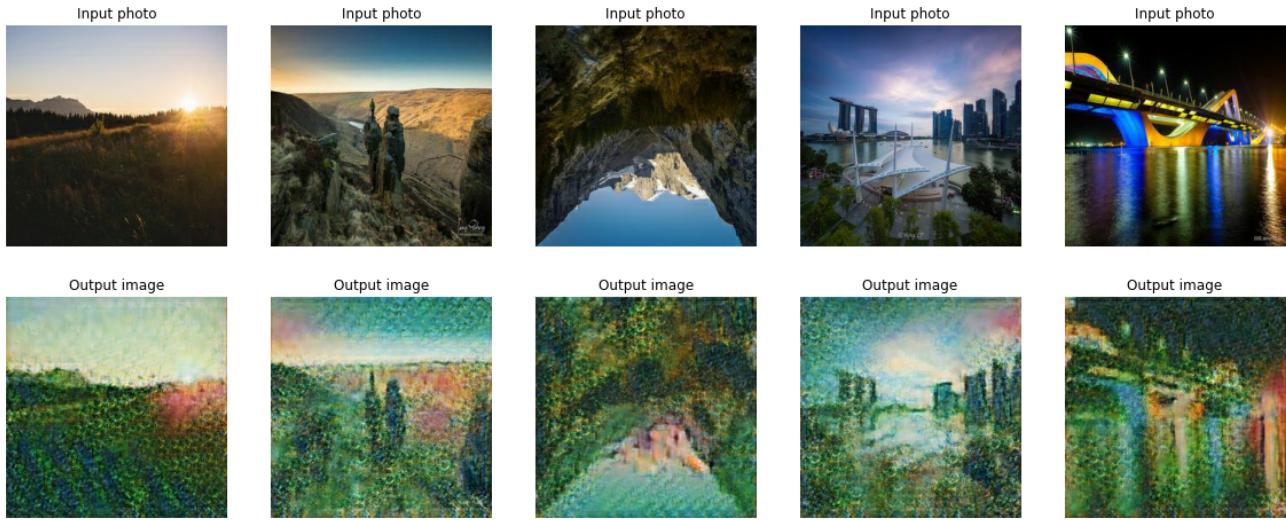


Epoch: 51/80	Generators Loss: 2.9565	Discriminators Loss: 0.1325
Epoch: 52/80	Generators Loss: 2.8789	Discriminators Loss: 0.1458
Epoch: 53/80	Generators Loss: 2.8637	Discriminators Loss: 0.1298
Epoch: 54/80	Generators Loss: 2.8644	Discriminators Loss: 0.1362
Epoch: 55/80	Generators Loss: 2.8138	Discriminators Loss: 0.1363
Epoch: 56/80	Generators Loss: 2.7655	Discriminators Loss: 0.1280
Epoch: 57/80	Generators Loss: 2.8331	Discriminators Loss: 0.1274
Epoch: 58/80	Generators Loss: 2.7315	Discriminators Loss: 0.1297
Epoch: 59/80	Generators Loss: 2.8161	Discriminators Loss: 0.1288
Epoch: 60/80	Generators Loss: 2.7957	Discriminators Loss: 0.1216



Epoch: 61/80 | Generators Loss: 2.7717 | Discriminators Loss: 0.1193

Epoch: 62/80	Generators Loss: 2.7705	Discriminators Loss: 0.1208
Epoch: 63/80	Generators Loss: 2.7081	Discriminators Loss: 0.1301
Epoch: 64/80	Generators Loss: 2.7171	Discriminators Loss: 0.1222
Epoch: 65/80	Generators Loss: 2.6859	Discriminators Loss: 0.1214
Epoch: 66/80	Generators Loss: 2.6967	Discriminators Loss: 0.1247
Epoch: 67/80	Generators Loss: 2.6797	Discriminators Loss: 0.1143
Epoch: 68/80	Generators Loss: 2.7060	Discriminators Loss: 0.1239
Epoch: 69/80	Generators Loss: 2.6109	Discriminators Loss: 0.1273
Epoch: 70/80	Generators Loss: 2.6857	Discriminators Loss: 0.1276



Epoch: 71/80	Generators Loss: 2.6193	Discriminators Loss: 0.1233
Epoch: 72/80	Generators Loss: 2.6080	Discriminators Loss: 0.1272
Epoch: 73/80	Generators Loss: 2.5859	Discriminators Loss: 0.1219
Epoch: 74/80	Generators Loss: 2.7110	Discriminators Loss: 0.1295
Epoch: 75/80	Generators Loss: 2.6496	Discriminators Loss: 0.1295
Epoch: 76/80	Generators Loss: 2.6237	Discriminators Loss: 0.1361
Epoch: 77/80	Generators Loss: 2.6987	Discriminators Loss: 0.1340
Epoch: 78/80	Generators Loss: 2.6473	Discriminators Loss: 0.1263
Epoch: 79/80	Generators Loss: 2.7237	Discriminators Loss: 0.1465
Epoch: 80/80	Generators Loss: 2.7680	Discriminators Loss: 0.1526

